

Enhanced Object Recognition Using CIFAR-10:

A comparative Study of TensorFlow, PyTorch and
Fine-Tuned AlexNet

Prepared By:

Maryam Ibrahim – 23012132

Mohamed Walied – 23011501

CIFAR-10 Dataset

airplane



automobile



bird



cat



deer



dog



frog



horse



ship



truck



Total 60,000 images

50,000 for training
(5000 images per class)

10,000 for testing
(1000 images per class)

Challenges in Dataset

- The brightness of some images was very low so we add in the augmentations some changes in the brightness with 20%.
- We also make in the augmentations a horizontal flip and random cropping to add more diversity on the dataset.
- therefore we resize it.

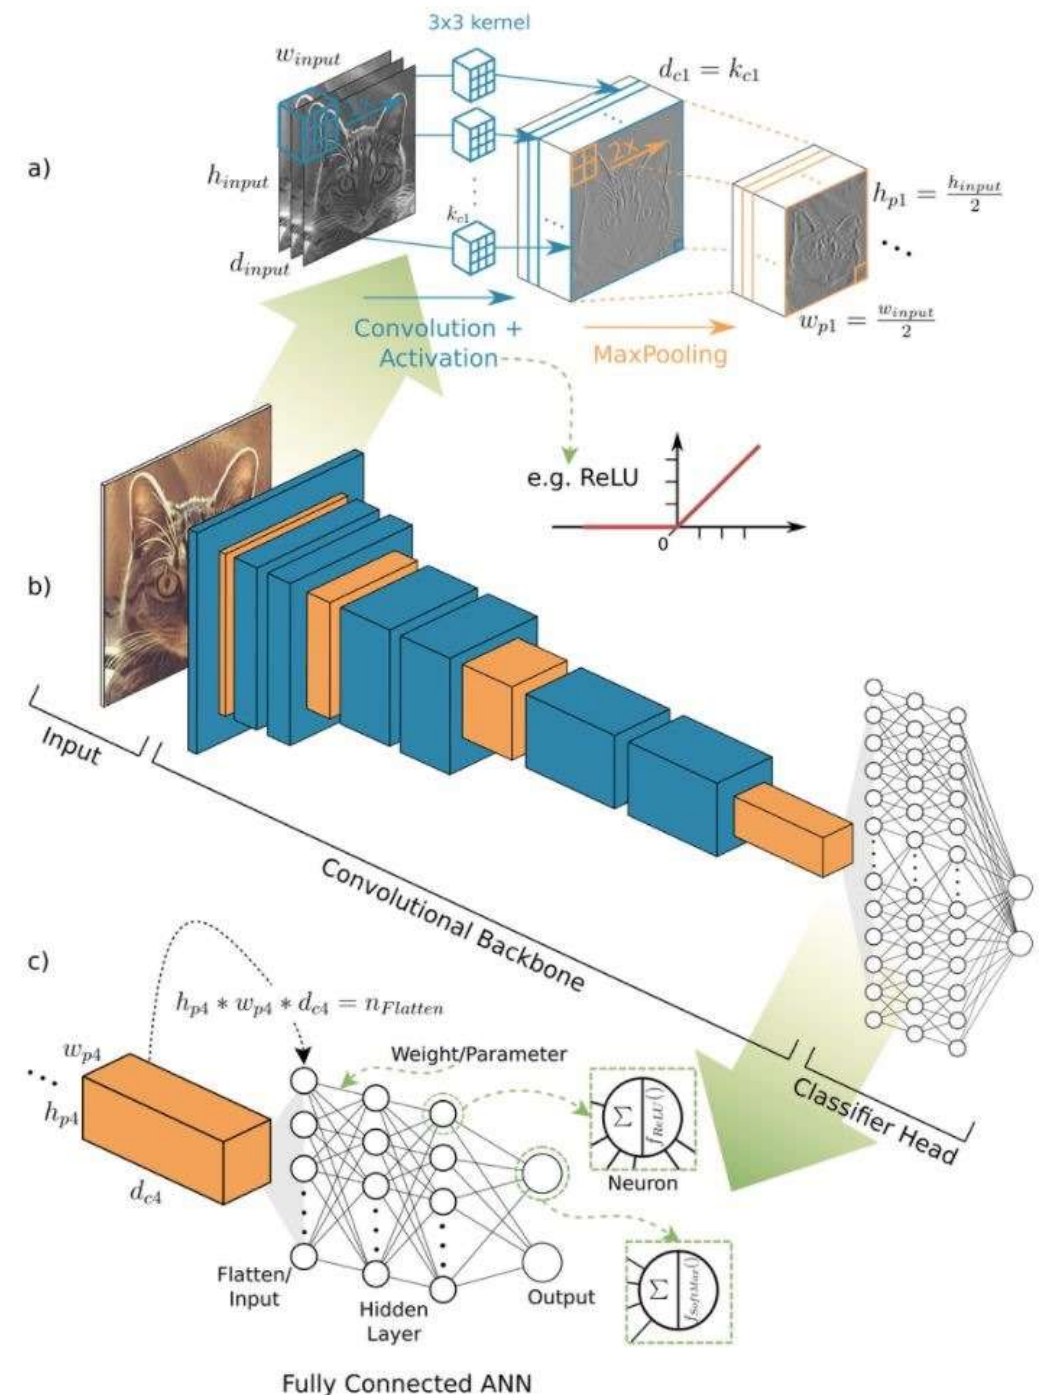
- Images should be normalized for good results in Neural Network.
- Images have to be transformed to 3D-Tensor so PyTorch can deal with it.
- AlexNet must take the data in dimensions of 224x224

CNN Models

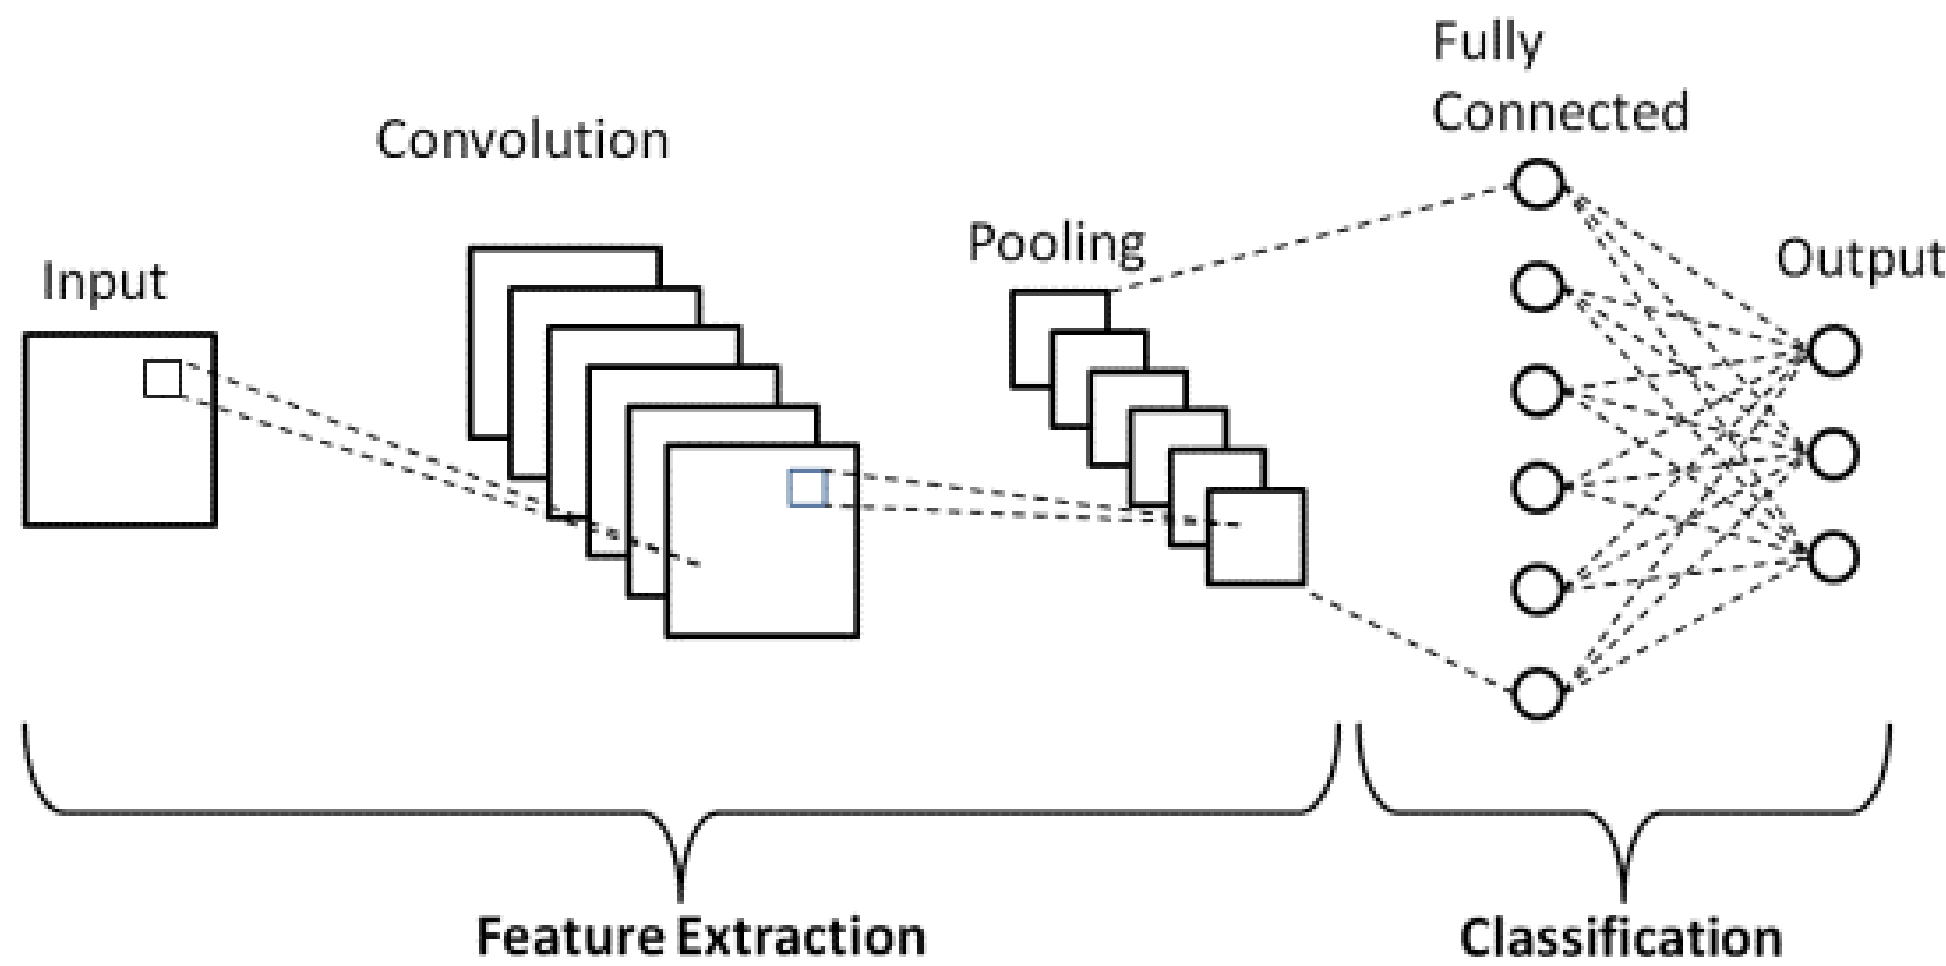
1. CNN architecture in PyTorch.

2. CNN architecture in TensorFlow.

3. Fine-Tuning AlexNet in PyTorch.



PyTorch CNN Model



Convolutional Operator

Input image

| | | | | |
|---|---|---|---|---|
| 9 | 4 | 1 | 2 | 2 |
| 1 | 1 | 1 | 0 | 4 |
| 1 | 2 | 1 | 0 | 6 |
| 1 | 0 | 0 | 2 | 2 |
| 9 | 6 | 7 | 4 | 1 |

Filter

| | | |
|---|---|---|
| 0 | 2 | 1 |
| 4 | 1 | 0 |
| 1 | 0 | 1 |


Output array

| | | |
|----|--|--|
| 16 | | |
| | | |
| | | |


$$\begin{aligned}\text{Output } [0][0] &= (9*0) + (4*2) + (1*4) \\ &+ (1*1) + (1*0) + (1*1) + (2*0) + (1*1) \\ &= 0 + 8 + 1 + 4 + 1 + 0 + 1 + 0 + 1 \\ &= 16\end{aligned}$$

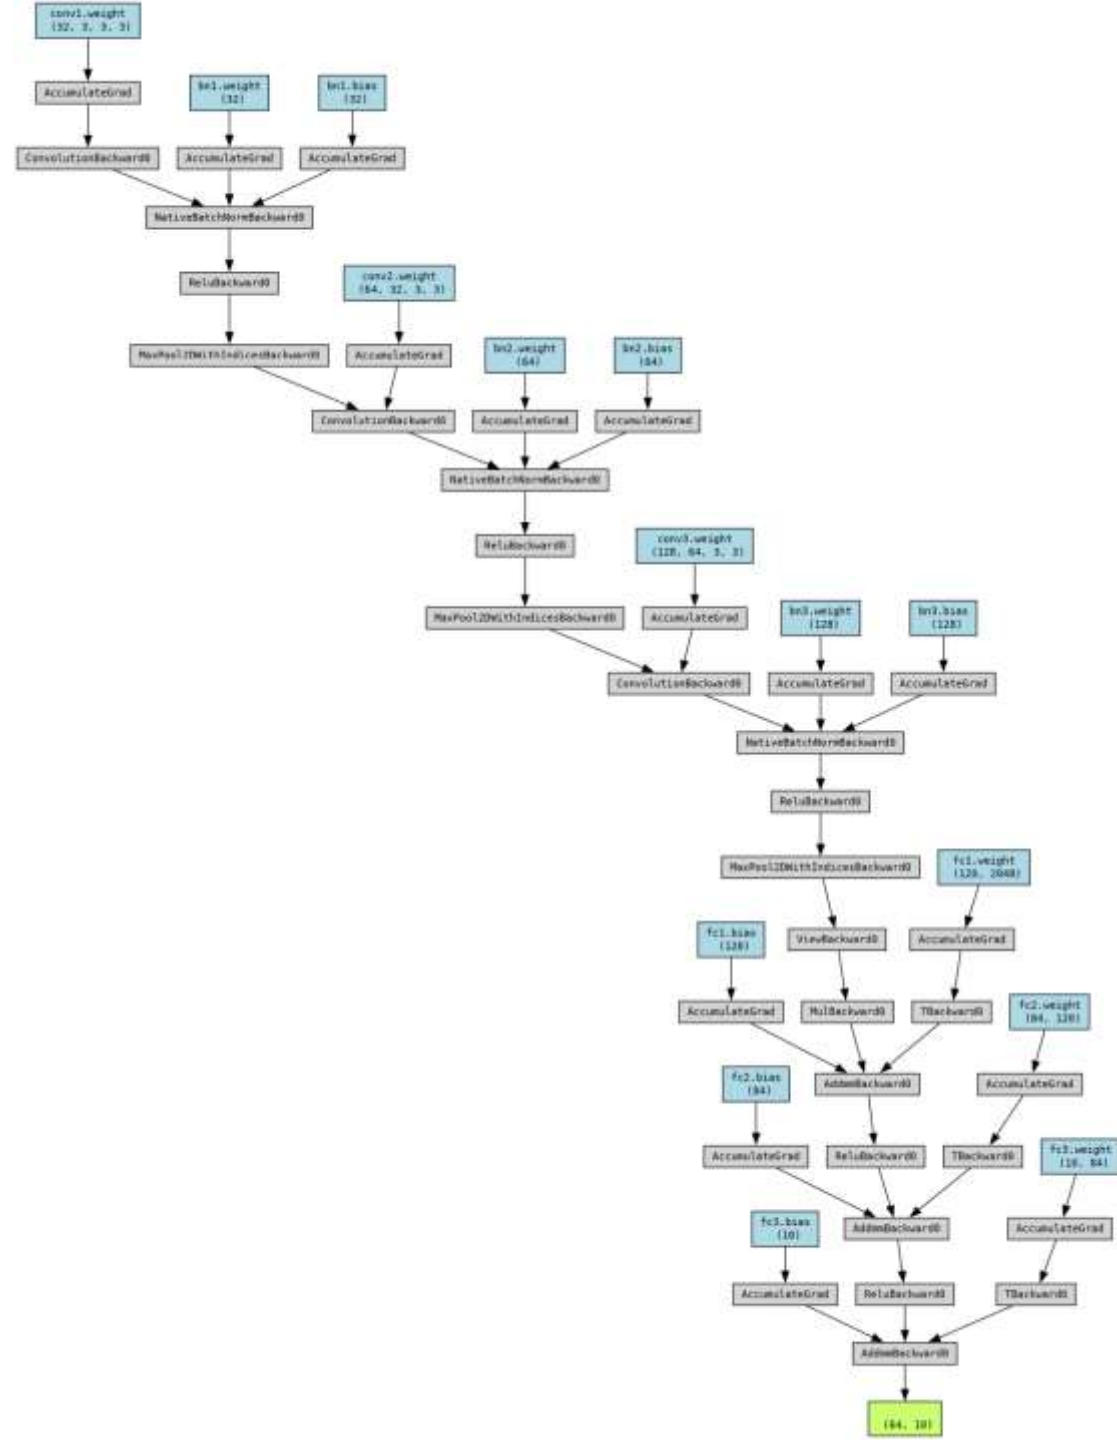
Convolutional Operator

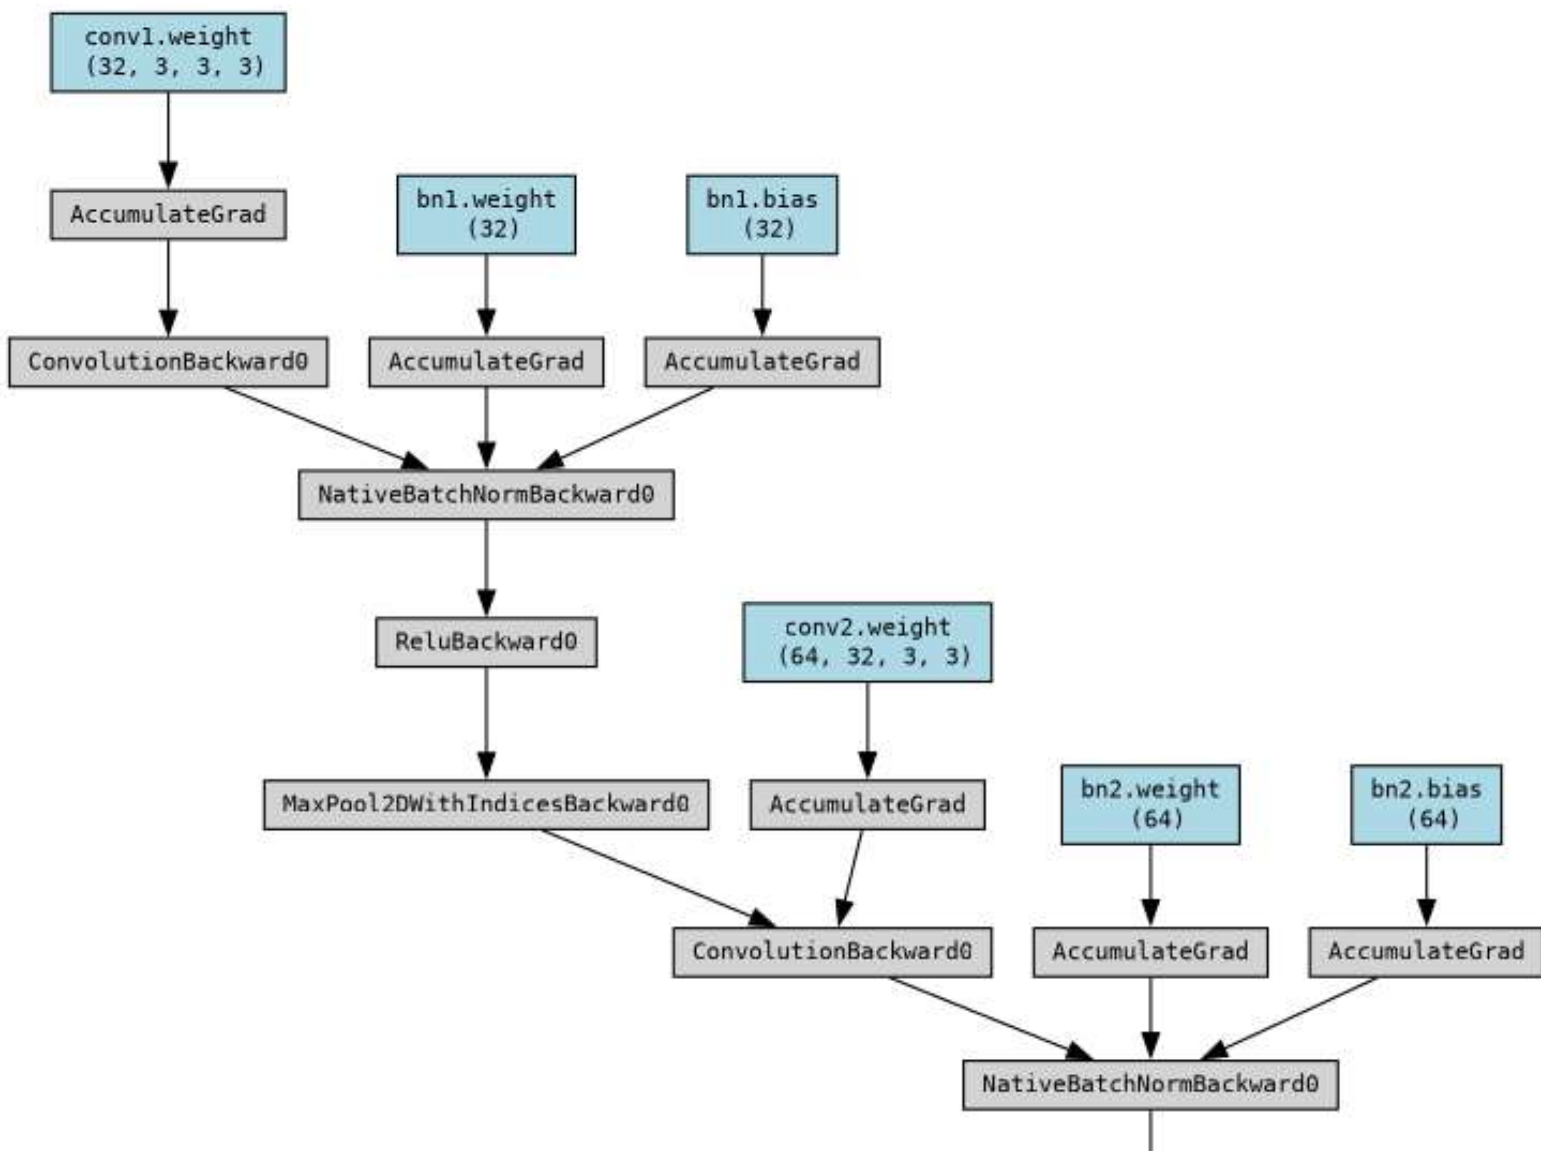
A Convolution Layer

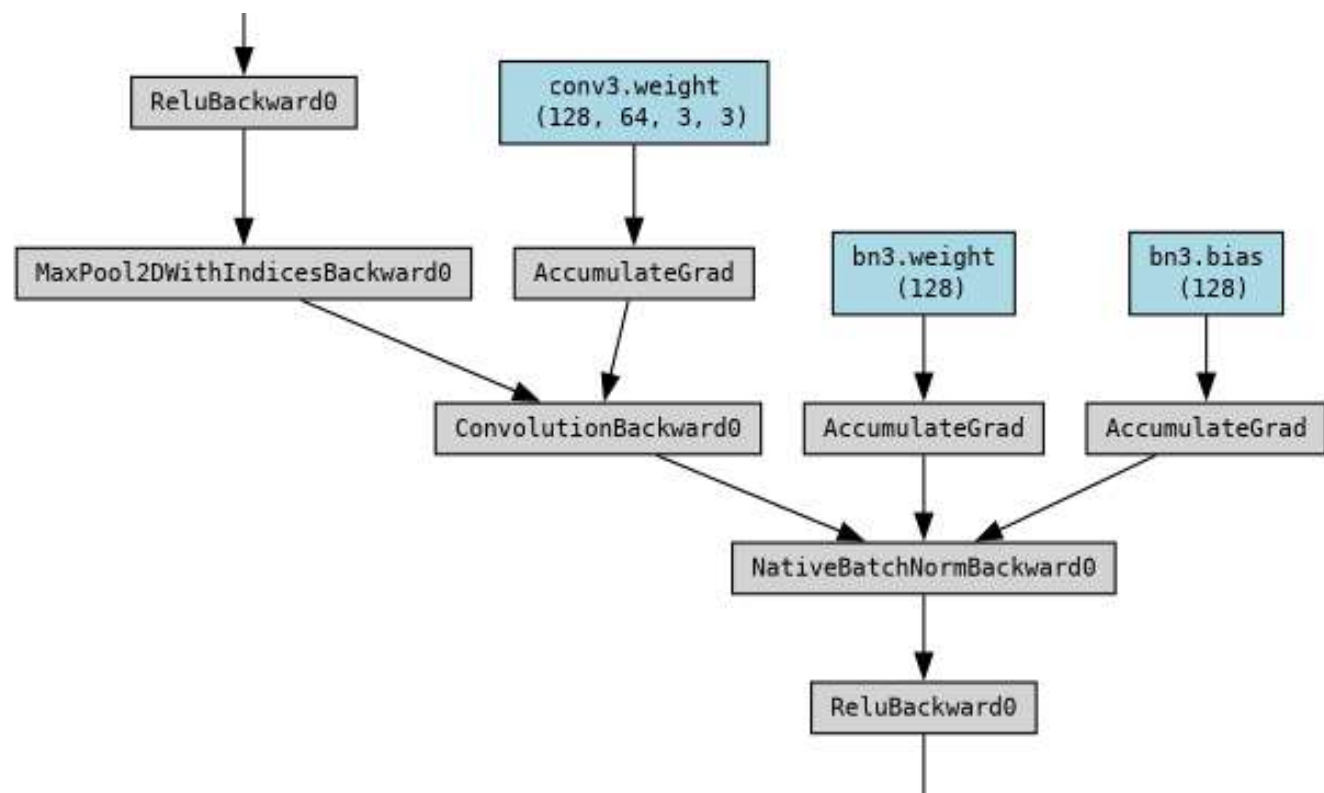


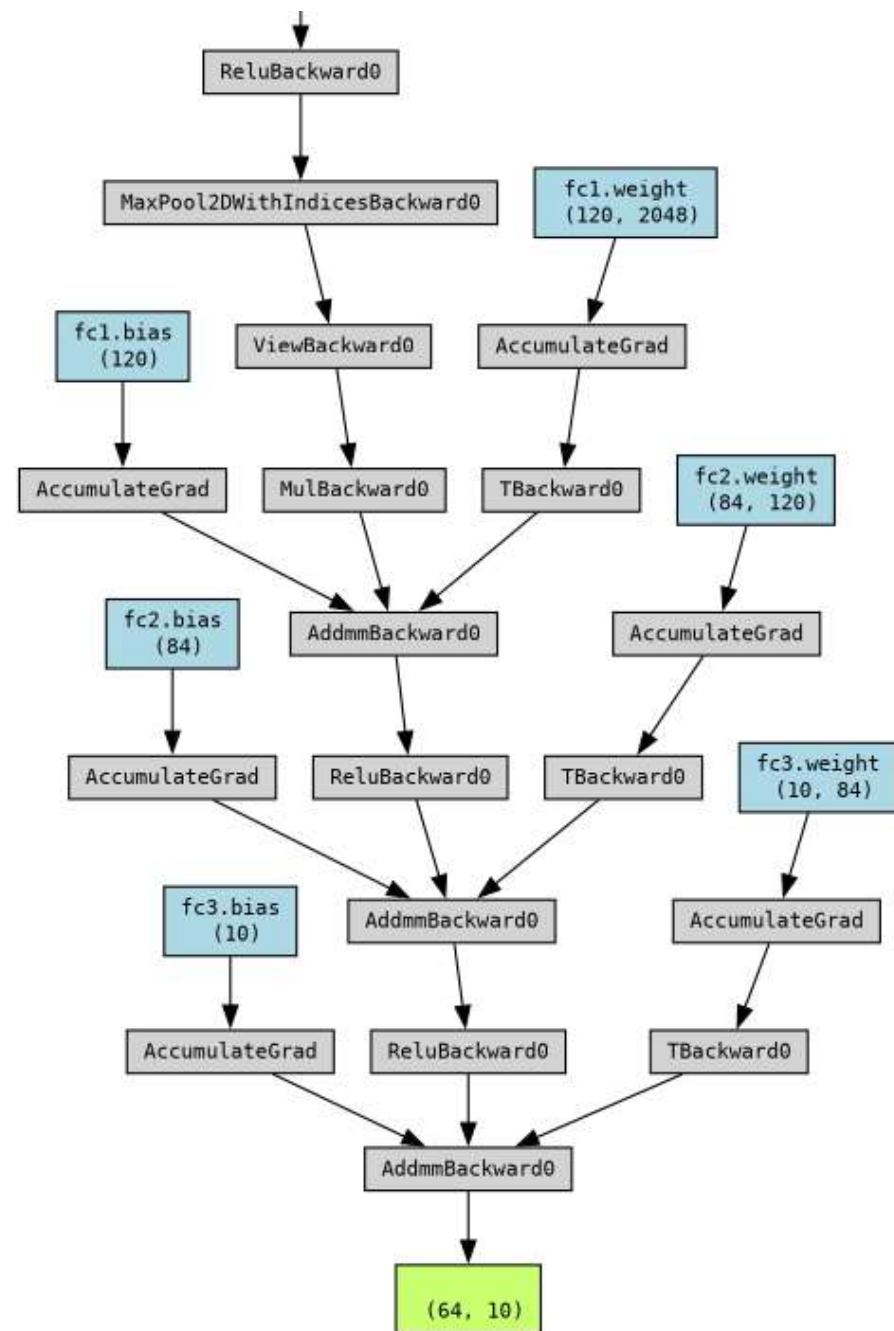
MaxPooling











```
class CNN(nn.Module):
    def __init__(self):
        super(CNN, self).__init__()
        # Convolutional and batchnorm layers
        self.conv1 = nn.Conv2d(3, 32, 3, padding='same', bias=False)
        self.bn1 = nn.BatchNorm2d(32)
        self.conv2 = nn.Conv2d(32, 64, 3, padding='same', bias=False)
        self.bn2 = nn.BatchNorm2d(64)
        self.conv3 = nn.Conv2d(64, 128, 3, padding='same', bias=False)
        self.bn3 = nn.BatchNorm2d(128)

        # Max pooling layers
        self.pool = nn.MaxPool2d(kernel_size=2, stride=2)


        # Calculate the output size after the convolutional layers
        self.fc1 = nn.Linear(128 * 4 * 4, 120) # Adjust input size accordingly
        self.fc2 = nn.Linear(120, 84)
        self.fc3 = nn.Linear(84, 10)

        # Adding Dropout layer to avoid underfitting
        self.dropout = nn.Dropout(0.5)


    def forward(self, x):

        x = self.pool(F.relu(self.bn1(self.conv1(x))))
        x = self.pool(F.relu(self.bn2(self.conv2(x))))
        x = self.pool(F.relu(self.bn3(self.conv3(x))))
        x = x.view(-1, 128 * 4 * 4) # Adjust input size accordingly
        x = self.dropout(x)
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = self.fc3(x)
        return x
```

Loss Curve (Training)



Training and Test Accuracy



Training dataset

Average loss:
0.4829

Accuracy: **82.86%**

Testing dataset

Average loss: **0.5722**
Accuracy: **80.64%**

Training time: **663.81 seconds**

No. of parameters: **350,342**

Epochs: **20**

Batchsize: **64**


Learning rate: **0.01**

Optimizer: **Adam**


Accelerator: **Tesla P100-PCIE-16GB (Kaggle)**

Some results

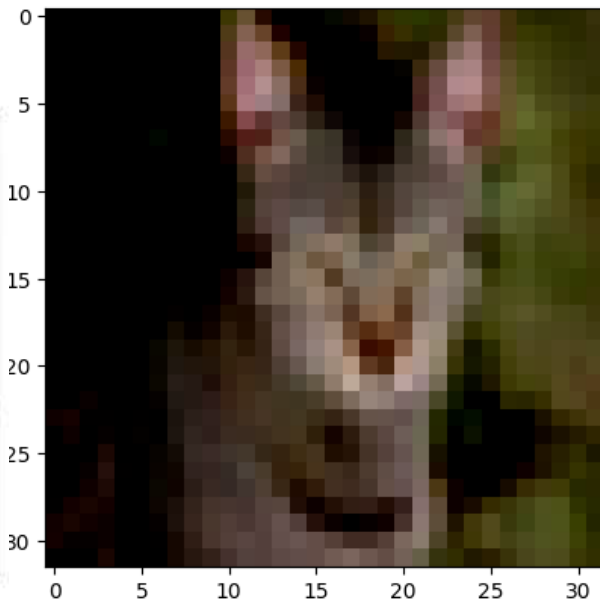
True Class: bird
Predicted Class: bird




True Class: airplane
Predicted Class: airplane



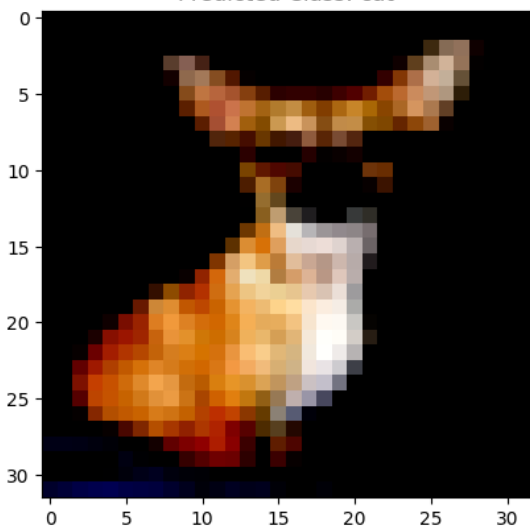
True Class: cat
Predicted Class: cat




True Class: cat
Predicted Class: cat




True Class: dog
Predicted Class: cat




True Class: dog
Predicted Class: cat



True Class: ship
Predicted Class: ship



True Class: frog
Predicted Class: frog



TensorFlow CNN Model

```
# Build a simple CNN
model = Sequential([
    Conv2D(32, (3, 3), activation='relu', input_shape=(32, 32, 3)),
    MaxPooling2D((2, 2)),
    Conv2D(64, (3, 3), activation='relu'),
    MaxPooling2D((2, 2)),
    Conv2D(128, (3, 3), activation='relu'),
    MaxPooling2D((2, 2)),
    Flatten(),
    Dense(64, activation='relu'),
    Dense(10, activation='softmax')
])
```

Training dataset

Average loss:

0.2692

Accuracy: **90.43%**

Testing dataset

Average loss: **1.2214**

Accuracy: **70.92%**

Training time: **84 seconds**

Epochs: **20**

No. of parameters: **224,842**

Batchsize: **32 (default)**

Learning rate: **0.01 (default)**

Optimizer: **Adam**

Accelerator: **Tesla P100-PCIE-16GB (Kaggle)**

Some

True: automobile, Pred: automobile



True: horse, Pred: horse



True: ship, Pred: ship




True: cat, Pred: cat



True: automobile, Pred: truck



True: dog, Pred: cat



True: frog, Pred: frog



True: cat, Pred: cat



AlexNet Fine-Tuning


ImageNet Classification with Deep Convolutional Neural Networks

Alex Krizhevsky
University of Toronto
kriz@cs.utoronto.ca


Ilya Sutskever
University of Toronto
ilya@cs.utoronto.ca

Geoffrey E. Hinton
University of Toronto
hinton@cs.utoronto.ca

2012



Fine-Tuning step



Pre-trained model

```
Sequential(
  (0): Conv2d(3, 64, kernel_size=(11, 11), stride=(4, 4), padding=(2, 2))
  (1): ReLU(inplace=True)
  (2): MaxPool2d(kernel_size=3, stride=2, padding=0, dilation=1, ceil_mode=False)
  (3): Conv2d(64, 192, kernel_size=(5, 5), stride=(1, 1), padding=(2, 2))
  (4): ReLU(inplace=True)
  (5): MaxPool2d(kernel_size=3, stride=2, padding=0, dilation=1, ceil_mode=False)
  (6): Conv2d(192, 384, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (7): ReLU(inplace=True)
  (8): Conv2d(384, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (9): ReLU(inplace=True)
  (10): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (11): ReLU(inplace=True)
  (12): MaxPool2d(kernel_size=3, stride=2, padding=0, dilation=1, ceil_mode=False)
)
-----
Sequential(
  (0): Dropout(p=0.5, inplace=False)
  (1): Linear(in_features=9216, out_features=4096, bias=True)
  (2): ReLU(inplace=True)
  (3): Dropout(p=0.5, inplace=False)
  (4): Linear(in_features=4096, out_features=4096, bias=True)
  (5): ReLU(inplace=True)
  (6): Linear(in_features=4096, out_features=1000, bias=True)
)
```


Fine-Tuned model

```
AlexNet(  
  (features): Sequential(  
    (0): Conv2d(3, 64, kernel_size=(11, 11), stride=(4, 4), padding=(2, 2))  
    (1): ReLU(inplace=True)  
    (2): MaxPool2d(kernel_size=3, stride=2, padding=0, dilation=1, ceil_mode=False)  
    (3): Conv2d(64, 192, kernel_size=(5, 5), stride=(1, 1), padding=(2, 2))  
    (4): ReLU(inplace=True)  
    (5): MaxPool2d(kernel_size=3, stride=2, padding=0, dilation=1, ceil_mode=False)  
    (6): Conv2d(192, 384, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (7): ReLU(inplace=True)  
    (8): Conv2d(384, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (9): ReLU(inplace=True)  
    (10): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (11): ReLU(inplace=True)  
    (12): MaxPool2d(kernel_size=3, stride=2, padding=0, dilation=1, ceil_mode=False)  
  )  
  (avgpool): AdaptiveAvgPool2d(output_size=(6, 6))  
  (classifier): Sequential(  
    (0): Dropout(p=0.5, inplace=False)  
    (1): Linear(in_features=9216, out_features=4096, bias=True)  
    (2): ReLU(inplace=True)  
    (3): Dropout(p=0.5, inplace=False)  
    (4): Linear(in_features=4096, out_features=1024, bias=True)  
    (5): ReLU(inplace=True)  
    (6): Linear(in_features=1024, out_features=10, bias=True)  
  )  
)
```

Training dataset

Average loss:
0.1290
Accuracy: **95.97%**

Testing dataset

Average loss: **0.5556**
Accuracy: **83.49%**

Training time: **740 seconds**

Epochs: **10**

No. of parameters: **60M (4,211,082 trainable)**

Batchsize: **64**

Learning rate: **0.01**

Optimizer: **Adam**

Accelerator: **Tesla P100-PCIE-16GB (Kaggle)**

Some results

True: cat
Pred: cat



True: airplane
Pred: airplane



True: cat
Pred: dog



True: ship
Pred: ship




True: dog
Pred: cat




True: bird
Pred: bird



True: frog
Pred: frog



True: bird
Pred: bird



Last Comparison

| Metric | Model 1 | Model 2 | Model 3 |
|-------------------|----------------------|----------------------|---------------------------|
| Training Accuracy | 82.86% | 90.43% | 95.97% |
| Testing Accuracy | 80.64% | 70.92% | 83.49% |
| Training Loss | 0.4829 | 0.2692 | 0.1290 |
| Testing Loss | 0.5722 | 1.2214 | 0.5556 |
| Training Time (s) | 663.81 | 84 | 740 |
| Parameters | 350,342 | 224,842 | 60M (4,211,082 trainable) |
| Epochs | 20 | 20 | 10 |
| Batch Size | 64 | 32 | 64 |
| Learning Rate | 0.01 | 0.01 (default) | 0.01 |
| Optimizer | Adam | Adam | Adam |
| Accelerator | Tesla P100-PCIE-16GB | Tesla P100-PCIE-16GB | Tesla P100-PCIE-16GB |

Thank You!